

MAXWELL J. YIN

maxwellyin@icloud.com | +1 (226) 998-0638 | Toronto, ON, Canada

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Ph.D. in NLP and Machine Learning Engineer specializing in LLM-powered product systems and evaluation. Experienced in designing and optimizing retrieval-augmented generation (RAG) systems, running large-scale experiments, and improving model quality, latency, and reliability in real-world applications.

EXPERIENCE

Huawei Technologies - Noah's Ark Lab (Canada)

08/2025 – 02/2026

Machine Learning Engineer

Toronto, ON

- **LLM Systems & Scaling:** Led 400M-parameter Transformer pretraining (40B tokens on 8×H100), delivering reproducible evaluation pipelines to support model selection and deployment decisions for downstream LLM applications (e.g., retrieval, reasoning, and user-facing AI systems) under compute constraints.
- **System Optimization & Deployment Readiness:** Profiled and resolved training/inference bottlenecks (GPU utilization, memory, latency). Optimized trade-offs to improve stability, throughput, and production readiness for downstream LLM applications.

The University of Western Ontario

09/2021 – 07/2025

Graduate Research Assistant

London, ON

- **Production Semantic Retrieval:** Built scalable retrieval system over 500K+ documents / 1.7M passages, reducing latency by 5× while improving ranking quality, enabling faster and more reliable responses in downstream LLM applications.
- **Evaluation & Experimentation:** Designed and ran large-scale experiments to evaluate retrieval and generation quality, including ablation studies and benchmarking pipelines, identifying failure modes and enabling data-driven improvements in model robustness and performance.

SELECTED PROJECTS

CineSeek: Production-Style Semantic Movie Search System

2026

[Live Demo](#) | [GitHub](#)

- **Retrieval-Augmented GenAI System:** Designed a low-latency semantic retrieval and ranking pipeline integrating query rewriting, FAISS-based ANN search, and LLM reranking, improving relevance while optimizing latency-quality trade-offs for real-world search applications.
- **Agentic Query Understanding:** Designed an agent layer with HTTP-based tool invocation, replacing subprocess-based execution with a unified MCP-style interface, enabling modular integration of retrieval and LLM components while reducing overhead and improving system reliability.
- **Full-Stack AI Deployment:** Built and deployed an end-to-end AI application with FastAPI backend, interactive React frontend, and real-time query serving, containerized with Docker and deployed over HTTPS behind a reverse proxy, with CI/CD automation via GitHub Actions for reproducible deployment and continuous delivery.

SKILLS

- **LLM Product & Evaluation:** RAG systems, embedding retrieval, A/B testing, offline/online evaluation, failure analysis, prompt engineering, model quality benchmarking.
- **LLM Systems & Training:** Pretraining, inference optimization (KV cache, batching, vLLM), post-training, transformer architectures, DeepSpeed ZeRO, mixed precision.
- **Production Infrastructure:** Python, PyTorch, HuggingFace Transformers, FastAPI, model serving, scalable pipelines, FAISS, Docker, React, LangChain, OpenAI API, SLURM, AWS, Git.

EDUCATION

The University of Western Ontario

09/2021 – 07/2025

Ph.D. in Computer Science (Natural Language Processing)

London, ON

- First-author publications in TACL, NAACL, and AAI ([Google Scholar](#)).

The University of Western Ontario

09/2019 – 08/2021

Master of Science in Computer Science (NLP & Machine Learning)

London, ON